

Practice 8 Memo of TCGI

Exercise 0.4

In order to configure both endpoints for the IPIP tunnel we run the following commands for R2:

```
R2:~# ip tunnel add tunnel0 mode ipip
local 198.51.100.2 remote 192.0.2.2 ttl
inherit nopmtudisc dev eth2
R2:~# ifconfig tunnel 1.2.3.4
R2:~# route add -net 192.168.0.0/24 dev
tunnel0
```

And for R1:

```
R1:~# ip tunnel add tunnel0 mode ipip
local 192.0.2.2 remote 198.51.100.2 ttl
inherit nopmtudisc dev eth2
R1:~# ifconfig tunnel0 4.3.2.1
R1:~# route add -net 172.16.1.0/24 dev
tunnel0
```

Exercise 0.5

While capturing in all SimNets we run `ping -c1 172.16.1.3` on `host2` in order to reach `host3`.

- The IP frames captured on `SimNet0` and `SimNet3` are the standard ones with no modification whatsoever. The `src @IP` show, is the private one, not the public.
- On the other hand, the frames captured on the other two `SimNets` have 2 IP headers. As it can be seen on the image:

* The source and destination `@IP`` of the outer header are the ones from the endpoints of the tunnel whereas the `@IP`` of the inner header feature the private `@IP`` of the end hosts.

TTL	SimNet0	SimNet1	SimNet2	SimNet3
Outer header	NA	63	62	NA
Inner header	64	63	63	62

The results of the TTL can be seen in the top table. Thanks to the option `inherit`, the outer header inherits the inner header TTL when entering the tunnel and the inner header decrements value for the TTL when exiting the tunnel. During transfer, only the outer TTL gets modified, the inner remains frozen.

In order to evaluate the value of the TTL, we are going to perform several pings increasing the field:

TTL=1

The packet gets dropped at the first hop, the R1. No outer header gets created because it is not transferred to the tunnel. The ICMP error message is the `time to live exceeded` generated by R1. We only see traffic on `SimNet0`.

TTL=2

As before, the ping did not reach the destination. It was lost on the RS router.

The ICMP message is the same as before but now it's generated by RS and only seen on `SimNet1`. This frame has as origin `192.0.2.1` and as destination `192.0.2.2`. The ICMP message does not get relayed backwards to the host (we don't see it on `SimNet0`) because the expired TTL belongs to the outer header. We do not see a `soft state` because no message has been sent from the encapsulator to the sender. We do see a behaviour compliant with the RFC 2003 because no messages have been sent to `host2`.

TTL=3

Now the response is received successfully. We do see the ICMP reply.

Exercise 6

Section 1

The MTU of the tunnel interfaces is 1480 bytes. We assume it's because the L2 protocol is ethernet, so the MTU is 1500 minus the outer header of 20 bytes. Now we restrict the MTU for `SimNet2` to 996 bytes.

After reestablishing the tunnel, we can see that the mtu on R1 is still 1480 but on the other side now has the expected value of 976, 20B less than the one set up for the link.

Section 2

So by running `ping -c1 -s500 -M want 172.16.1.3` we set a payload length of 500 bytes, and 1 to the `don't fragment` flag. This option only allows fragmentation locally, not in path.

We can see two ICMP frames on each `SimNet`, one for the request and the other one for the reply. The sizes on the private networks are 542B whereas on the tunnel, thanks to the outer header, 562B.

The packets have not been fragmented because they do not exceed the minimum MTU (996B) and as expected the DF flag is set to 1 on the outer and inner headers of the request.

Section 3

Now we allow fragmentation in path: `ping -c1 -s500 -M dont 172.16.1.3`.

The results are identical with the DF flag set to 0 in all headers.

Section 4

Now we execute: `ping -c1 -s500 -M do 172.16.1.3`. This option prohibits all fragmentation, even local one. Again, no difference.

In order to have the maximum size allowed, we need to send a ping with parameter `-s948` which results from 996-20-8.

The -M want option

After clearing the cache, we run `host2:~# ping -c1 -s1000 -M want 172.16.1.3`. The ping is not successful.

With `want`, we only deactivate DF when needed.

We can see the first ICMP frame in `SimNet0`, then inside `SimNet1`, it gets dropped and notified to `R1` with an ICMP frame due to its size.

The lengths are 542 in the private nets and 562 in the transport ones. This length does exceed the minimum MTU.

All IP headers except from the ICMP error one have DF set to 1 (inner and outer).

The error is originated on `R1` and notified to `R1`. The notified mtu is 996.

In my humble opinion, the destination should be `host2`. It is not acting as expected because an error code 4 should be notified to the sender.

Soft State

After running `ping -c2 -s1000 -M want -i1 172.16.1.3` we can see the following output:

```
host2:~# ping -c2 -s1000 -M want -i1 172.16.1.3
PING 172.16.1.3 (172.16.1.3) 1000(1028) bytes of data.
From 192.168.0.1 icmp_seq=2 Frag needed and DF set (mtu = 976)
```

```
--- 172.16.1.3 ping statistics ---
2 packets transmitted, 0 received, +1 errors, 100% packet loss, time 1009ms
```

We can now see an error pop up. If we analyze the captured traffic, we can observe that the first ICMP ping request message gets discarded on `RS`. Is this router who notifies then `R1` as before. The difference comes with the second frame, this one gets dropped at `R1` and is this router who notifies `host2` that the length is exceeding the MTU. We can't see the second frame on `SimNet2`.

The DF is set on all ICMP ping request frames. We can now conclude that `R1` maintains a soft state of the MTU. We can retrieve it by running `ip route show cache`:

```
local 192.0.2.2 from 192.0.2.1 dev lo
src 192.0.2.2
    cache <local,src-direct> iif eth2
192.168.0.2 dev eth1 src 192.168.0.1
    cache ipid 0x0b79
172.16.1.3 dev tunnel0 src 192.0.2.2
    cache expires 179sec ipid 0xc69d
mtu 976
198.51.100.2 from 192.0.2.2 via
192.0.2.1 dev eth2
    cache expires 178sec ipid 0x107c
mtu 996
```

It reads: in order to arrive to `@IP1` from `@IP2` via `@IP3`, this restriction is applied. It has an expiration time.

Now, by increasing the count to three, we can see that the third try succeeds.

By analysing the captured traffic, we can see that the third ping does not have the DF activated and comes fragmented from `host2` (DF activated on in the first piece).

SimNet0

The sizes of the fragments are 986 and 90 each. (With the ethernet header)

SimNet1

The sizes of the fragments are 1010 and 110 each (with the ethernet header). Provided that the ethernet header size is 14B, neither of these frames exceed the minimum MTU.

The fragmentations in `SimNet1` and `SimNet0` are different because in `SimNet1` a 20 byte header must be added. The new outer header must have track of both fragments with the fragment offset field (which now keeps also track of the inner header). In order to have data sizes multiple of 8, four bytes from the second fragment are moved to the first one. This can be done because DF is disabled.

SimNet3

We see the same ICMP requests as in `SimNet2`.

All the fragmented fragments have DF disabled-

The ICMP response fragments fit natively the MTU of the tunnel.

Finally, we run `ping -c2 -s 1460 -M want -i1 172.16.1.3`:

The second request is the one that succeeds. Because it exceeds the MTU of the first link, the `R1` generated the error right away so the next message gets fragmented from the `host2` for a mtu of 1480. This fragmentation puts the DF to 0 so on `RS` the fragments get resized to fit the channel. The response fits the minimum MTU natively.

The -M do option

Now we run `ping -c3 -s 1000 -M do -i1 172.16.1.3`. It does not succeed. The packets are sent as:

1. Does get sent and dropped at `RC` and notified to `R1` not to `host2`
2. Does get sent and dropped at `R1` and notified to `host2` from `R1` (soft state)
3. Does not get sent at all.

So with `do` does not send if he knows that will not succeed.

The -M dont

We set the `dont` so we will set always the DF to 0. We run `ping -c3 -s 1000 -M dont -i1 172.16.1.3`. All packets succeed because they have the

DF flag unset. They get fragmented at the entrance of the 996 mtu link.

The fragmentation in path has less fragmentations needed. It also takes advantage of the full MTU of the first link.

Exercise 7

The -M want

After setting up the scenario, we run the want ping:

1. The first packet gets dropped at RC and notified to R1 silently.
2. The second one gets dropped at R1 and notified to host2 with the same mtu as in the exercise before.
3. The third frame gets fragmented and not forwarded to RC because pmtudisc does not allow fragmented frames to the tunnel. Its purpose is to avoid fragmentation. It gets silently discarded.

The -M do

After setting up the scenario we run the do ping:

1. The first packet gets dropped at RC as before.
2. The second one behaves as the second one from before.
3. The third one does not get even sent by ping because it realizes it will get dropped by R1.

The outer header of the first packet at SimNet1 has DF set to 1 (due to the pmtu disc) whereas the inner has it set to 0 (due to the dont on ping).

The -M dont

After setting up the scenario we run the dont ping:

1. The first packet gets dropped and notified to R1 at RC.
2. The frame is silently discarded at R1.
3. The frame is silently discarded at R1.

They are silently discarded because the notification would be fragmentation needed and the frames are set DF so they can be fragmented.

So the pmtudisc does not allow fragmented frames or frames that would be fragmented in the tunnel

Exercise 0.8

The sender will all always use the minimum MSS.

Section 0.8.1

We start the connection between the host2 and host3. We can see the SYN, SYN/ACK and the final ACK of the three way handshake. The MSS advised at SYN from host2 is 1460. On the other hand the MSS advised by host3 by SYN/ACK is also 1460 because host3 does not see the tunnel MTU.

Section 0.8.2

Now we are going to try to send a file.

- In the three way handshake the MSS was 1460.

- The first data packages from the server are refused by the router due to the mtu. The error is generated on the R2 and propagated to the host3.
- host3 is in charge of solving the problem by reducing the frame size.
- The MSS has not changed because it is a local value.
- The file is transmitted properly.

Now we do the reverse operation:

```
host2:~# cat /etc/services | nc 192.168.0.2 12345
```

- Now we can see the same error than before but now on SimNet0
- The first ICMP error is generated by NS and reported to R1. Then the host sends 3 more frames that will also get discarded but with an ICMP error message from R1.
- The file was transmitted correctly and the TCP client is host2

Section 0.8.3

- The three way handshake does work because the frames are less than 100B long.
- Yes we have a problem. We see an ICMP error at SimNet0 reported by R1 because the first data messages exceed the 1480 mtu (1500).
- The next messages silently die at RS because the R1 router does not accept input ICMP messages. The file cannot be transmitted.
- If we use the forward chain it would work, because we are not relaying ICMP messages we are using ICMP.

Section 0.8.4

The MSS should be $996 - 20 - 20 - 20 = 936$. We change it on the host3 side.

Change the MSS in the host

We run the command `host3:~# ip route add 192.168.0.0/24 via 172.16.1.1 advmss 936` And now everything works.

Change the MSS in the router

In order to set up the scenario we run on the R2 `iptables -t mangle -A FORWARD -o tunnel0 -p tcp --syn -j TCPMSS --set-mss 936`.

We use the --syn instead of what we did before due to the direction and the profiles of the hosts. (client and server)

Practice 7 memo of TCGI

Exercise 1

Section 1

After running the default configuration command we can see that the configuration of the mentioned hosts we configure the filtering tables of host1. The rules are:

- Block all entering ICMP traffic.

We run on host1: `iptables -t filter -A INPUT -p ICMP -j DROP`

And now we test the configuration by *pinging* the host1 from Rint.

- We can capture the echo-request generated by Rint.
- We can't not capture the echo-reply because the echo-request packet gets dropped when it arrives.
- It gets silently dropped

Now we perform the reverse operation.

- We can capture both the echo-request and the echo-reply but the host1 reports that the package has been lost.
- The cause of this behaviour is the filter. The package is received but silently drop and never gets served to the application.

Section 2

The next step is to remove the before applied filter by running: `iptables -t filter -D INPUT -p ICMP -j DROP` And we test the configuration with a simple ping. It's time to configure by:

- The host1 must be able to ping Rint.
- The host1 must no reply any ping request.

In order to configure this set-up we run: `iptables -t filter -A INPUT -p ICMP --icmp-type echo-request -j DROP`.

- When host1 pings Rint we see the normal behaviour.
- When Rint pings hosts we see the same procedure than before when all ICMP packets where dropped.

Section 3

Now we configure the Rint router as the *network guard*. We first delete all entries of the filtering tables of host1 by running `iptables -t filter -D INPUT -p ICMP --icmp-type echo-request -j DROP`. After checking the connection with Net1 via ping ping we can proceed by setting up the requested escenario with the following commands:

```
Rint:~# iptables -t filter -A FORWARD -p ICMP -s 192.168.1.0/24 --icmp-type echo-request -j ACCEPT
Rint:~# iptables -t filter -A FORWARD -p ICMP --icmp-type echo-request -j DROP
```

Now the ping from host1 to www works but not vice-versa.

In order to setup the TCP filtering, we run the following commands:

```
Rint:~# iptables -t filter -A FORWARD -s 192.168.1.0/24 -p tcp --syn -j ACCEPT
Rint:~# iptables -t filter -A FORWARD -p tcp --syn -j DROP
```

Now try a nc from host1 to www and it works but not vice-versa.

The last configuration is the UDP, we add:

```
-A FORWARD -s 172.16.1.5/32 -p udp -j ACCEPT
-A FORWARD -d 172.16.1.5/32 -p udp -j ACCEPT
-A FORWARD -p udp -j DROP
```

We try this configuration with the command `dig dns.practnet.tcgi`

This is the final configuration file:

```
Rint:~# cat tables.txt
# Generated by iptables-save v1.4.2 on Sun Apr 28 16:11:18 2019
*filter
:INPUT ACCEPT [4:514]
:FORWARD ACCEPT [5:396]
:OUTPUT ACCEPT [4:278]
-A FORWARD -s 192.168.1.0/24 -p tcp -m tcp --tcp-flags FIN,SYN,RST,ACK SYN -j ACCEPT
-A FORWARD -p tcp -m tcp --tcp-flags FIN,SYN,RST,ACK SYN -j DROP

-A FORWARD -s 192.168.1.0/24 -p icmp -m icmp --icmp-type 8 -j ACCEPT
-A FORWARD -p icmp -m icmp --icmp-type 8 -j DROP

-A FORWARD -s 172.16.1.5/32 -p udp -j ACCEPT
-A FORWARD -d 172.16.1.5/32 -p udp -j ACCEPT
-A FORWARD -p udp -j DROP
COMMIT
# Completed on Sun Apr 28 16:11:18 2019
```

Exercise 2

The first step is to try to ping the test machine from www. If we capter the traffic on SimNet0 we can see that the destiation @IP is correct but the source @IP is private, not 'public'. So this is why we do not receive any reply.

After running Rbcn:~# `iptables -t nat -A POSTROUTING -o eth2 -j SNAT --to 10.0.2.2`, the pting command works like a charm.

Now we want to give test acces to the Net1. If we run the ping from test we get the Network unreachable error. In order to be able to acces the net, we run Rbcn:~# `iptables -t nat -A PREROUTING -p tcp --destination-port 80 -j DNAT --to 172.16.1.2`